

Second Life AI — Technical Specification

Six-class waste image classification · NYU Shanghai · 22 July 2026 · Answers to the technical questionnaire

READ THIS FIRST — WHICH MODEL THESE NUMBERS DESCRIBE

The project has two phases. **Every measured number in this document comes from Phase 1:** a ResNet-50 trained on the 2,527-image TrashNet corpus. That is the model currently deployed in the website and the only one that has been trained.

Phase 2 is a larger five-source corpus under a 30M-parameter competition cap. Its data pipeline is built and verified, but **it has not been trained yet**, so no accuracy figure exists for it. Where Phase 2 appears below it is labelled as planned, not measured.

I. Dataset and Preprocessing

I.1 Train / test split ratio and image counts

There is a validation set, and it is separate from the test set. We use **three** partitions, not two, and the split is made **by near-duplicate group rather than by image** — see the note below, because this changes what the ratio means.

Table 1. Phase 1 — TrashNet, the trained model. Verified in `reports/locked_baseline/split_arithmetic.json`.

PARTITION	IMAGES	SHARE	PURPOSE
Fold-0 training	1,732	68.5%	Gradient updates
Fold-0 validation	434	17.2%	Model selection, early stopping
Test (quarantined)	361	14.3%	Used exactly once
Total inventory	2,527	100%	

The 2,166 non-test images carry a fold label 0–4, so the same data supports five-fold cross-validation; validation fold sizes are 434, 433, 433, 433 and 433. All reported single-run results use fold 0 as validation.

WHY THIS IS NOT A PLAIN 80/20 SPLIT

TrashNet contains many photographs of the same physical object. A random image-level split can put one view of a bottle in training and another view in validation, so the model scores by recognising the object rather than the material. We therefore detect near-duplicates — 64-bit perceptual hash at Hamming distance ≤ 6 , plus ResNet-18 embedding cosine similarity ≥ 0.96 — merge them into groups with union-find, and assign *groups* to partitions using `StratifiedGroupKFold`. No group spans two partitions, which is asserted in code at every load.

Table 2. Phase 2 — unified corpus. Built and verified; **not yet trained on**. The training machine downloads five sources and rebuilds, so the final totals will be larger.

PARTITION	IMAGES	SHARE
Train + validation (5 folds of 1,912)	9,560	84.6%
Test (quarantined)	1,739	15.4%
Total after filtering	11,299	100%

I.2 Preprocessing and data augmentation

Augmentation is applied **only** to training images, and **only after** the split is decided — it lives inside the training dataset object, so an evaluation dataset structurally cannot receive it. Implemented in `model/src/data.py`.

Table 3. Training transform, `aug: basic` — the configuration used by the trained model. Order is exactly as applied.

#	OPERATION	PARAMETERS
1	RandomResizedCrop	size 224, scale (0.6, 1.0)
2	RandomHorizontalFlip	p = 0.5
3	RandomRotation	$\pm 15^\circ$
4	ColorJitter	brightness 0.2, contrast 0.2, saturation 0.2, hue 0.05
5	ToTensor	—
6	Normalize	ImageNet mean (0.485, 0.456, 0.406), std (0.229, 0.224, 0.225)
7	RandomErasing	p = 0.0 in this config — present in code but inactive

Table 4. Validation and test transform — deterministic, no randomness.

#	OPERATION	PARAMETERS
1	Resize	255 (= 224 × 1.14, the standard 256/224 crop ratio)
2	CenterCrop	224
3	ToTensor	—
4	Normalize	ImageNet mean / std, as above

Test-time augmentation (adopted). At inference the softmax is averaged over four deterministic views: (1) the standard eval transform; (2) the same with a horizontal flip; (3) the full image resized to 224×224 with no crop; (4) resize to 302 then centre-crop 224. This is the single change that passed our acceptance gate, worth +1.84 percentage points.

Other augmentation modes implemented but not used by the trained model: `aug: modern` replaces steps 1–4 with `RandomResizedCrop` + `RandomHorizontalFlip` + `TrivialAugmentWide`; Mixup and CutMix are wired up but set to 0.0; a per-class "boost" transform for minority classes exists and is disabled by default. The Phase 2 configs enable `modern`, mixup 0.2, cutmix 0.2 and random erasing 0.25, on the reasoning that these were marginal at 2.5k images and should pay off at four times the data.

II. Model Architecture and Training Parameters

II.1 Final network architecture

ResNet-50, loaded through `timm` as `resnet50.tv2_in1k` (the torchvision IMAGENET1K_V2 weights). The 1000-class ImageNet classifier is replaced with a fresh 6-class linear layer; the convolutional trunk is unchanged. Measured parameter count: **23.5M**, which is under the competition's 30M cap.

The intended final model is ConvNeXtV2 Tiny (`convnextv2_tiny.fcmae_ft_in22k_in1k`, 27.9M parameters). Three alternatives are configured and verified to build — CAFormer-S18 (24.3M), DINOv2 ViT-S/14 (22.1M), EfficientNetV2-S (20.2M) — but we do not plan to run them; see §V.4 for the reasoning. The cap is enforced in code: `build_model` counts parameters and raises before training starts, so an ineligible configuration fails in the first second rather than after hours.

ConvNeXtV2 Tiny is chosen as the *safest* option rather than the theoretically best one. It has strong pretrained features from FCMAE pre-training plus ImageNet-22k fine-tuning, it sits comfortably under the 30M cap with headroom, and it is a well-understood convolutional architecture that behaves predictably on a fixed budget. With only 13 hours of rented GPU time, reliability is worth more than the chance of a marginally better architecture.

II.2 Transfer learning: what was frozen, what was fine-tuned

Yes — ImageNet pre-trained weights, fine-tuned in two phases. **No layer is permanently frozen.**

PHASE	EPOCHS	TRAINABLE	LEARNING RATE
1 — head only	5	Backbone frozen (<code>requires_grad = False</code> on every non-classifier parameter); only the new 6-class head trains	head 1×10^{-3} , constant
2 — full fine-tune	40	Entire network unfrozen, in two parameter groups	backbone 1×10^{-5} , head 1×10^{-3} (100× ratio)

The head phase exists so the randomly initialised classifier does not send large, meaningless gradients back through pre-trained features on the first steps. The 100× discriminative learning-rate ratio in phase 2 is deliberate: in our experience on this dataset, too high a backbone learning rate is the single most common cause of a collapse from ~93% to ~75%.

II.3 Core hyperparameters

Table 5. Complete configuration of the trained model. Source: `configs/baseline.yaml`, reproduced in the `config_json` column of `experiments.csv` under hash `ab586f1a72`.

SETTING	VALUE
Optimizer	AdamW
Initial learning rate	Head 1×10^{-3} ; backbone 1×10^{-5} (separate parameter groups)
LR scheduler	LambdaLR implementing 3-epoch linear warmup, then cosine decay across the 40 fine-tune epochs. Constant ($\times 1.0$) during the 5 head epochs.
Loss function	CrossEntropyLoss with <code>label_smoothing = 0.1</code> . No class weights in the loss — imbalance is handled by the sampler instead (see below). Validation loss uses plain CrossEntropyLoss with no smoothing, so it stays comparable across configs.
Batch size	32
Total epochs	45 (5 head + 40 fine-tune). All 45 ran; early stopping was never triggered.
Weight decay	0.05, applied to both parameter groups
Dropout	None added. ResNet-50 has no dropout layer, and we did not insert one — regularization comes from weight decay, label smoothing, augmentation and EMA.
Class imbalance	WeightedRandomSampler with inverse-frequency weights, so the 137-image <i>trash</i> class is replayed roughly 4x per epoch
EMA	Enabled, decay 0.9998 (see the note below — it did not help)
Early stopping	Patience 10 epochs on validation accuracy (not reached)
Mixed precision	Enabled (AMP)
Random seed	42, set across Python, NumPy and PyTorch

A RESULT WORTH REPORTING: EMA HURT

We kept an exponential moving average of the weights, which is normally a free improvement. Here it was not. The EMA weights reached **88.94%** validation accuracy while the raw weights from the same run reached **93.09%** — a 4.15-point gap in the wrong direction, most likely because the decay of 0.9998 averages over a horizon far longer than a 45-epoch run on 1,732 images. The adopted checkpoint is the **raw** one. This is the kind of default that is worth checking rather than assuming.

III. Experimental Results and Evaluation

III.1 Accuracy on the independent test set

SINGLE-USE MEASUREMENT — THIS CAPTION MUST TRAVEL WITH THE NUMBER

322 / 361 = 89.20% on the held-out TrashNet test set. Measured **once**, before the later tuning experiments, and never used for model selection. It is **not** the score of the final optimized model, and it cannot be re-measured — the test set is spent, and a guard file now blocks a second evaluation programmatically.

Table 6. All measured results, fold-o validation (434 images) unless stated.

CONFIGURATION	ACCURACY	CORRECT	MACRO-F1	STATUS
ResNet-50, single view	93.09%	404/434	0.9196	Starting point
ResNet-50 + 4-view TTA	94.93%	412/434	0.9414	Adopted
Progressive resize 384px + TTA	94.93%	412/434	0.9442	Rejected
Modern augmentation + TTA	95.62%	415/434	0.9421	Rejected
Weight decay 0.10 + TTA	94.70%	411/434	0.9390	Rejected
Test set, adopted model	89.20%	322/361	—	Spent

Three configurations that raised or held accuracy were rejected. We fixed an acceptance rule *before* running them: a change is adopted only if it gains at least seven validation images, loses no macro-F1, and costs no class more than five points of recall. Modern augmentation had the highest raw accuracy we ever measured and still failed, because *trash* recall fell from 20/23 to 18/23 — the model got worse at the class it was already worst at.

III.2 Confusion matrix, precision, recall and F1

Table 7. Per-class metrics for the adopted configuration (ResNet-50 + 4-view TTA, fold-o validation). Source:

reports/baseline_resnet50/report_baseline_resnet50_val_fold0_tta.txt .

CLASS	PRECISION	RECALL	F1	SUPPORT
cardboard	1.0000	0.9275	0.9624	69
glass	0.9750	0.9070	0.9398	86
metal	0.9437	0.9437	0.9437	71
paper	0.9528	0.9902	0.9712	102
plastic	0.9011	0.9880	0.9425	83
trash	0.9091	0.8696	0.8889	23
Macro average	0.9469	0.9376	0.9414	434
Weighted average	0.9510	0.9493	0.9492	434




Figure 1. Confusion matrix, fold-0 validation. The dominant error is **glass** → **plastic** (4 images) with **glass** → **metal** (3) next — transparent and reflective surfaces, which is the expected failure mode. *trash* is the weakest class at 87% recall on only 23 images, so each error there moves its recall by 4.3 points.

III.3 Training loss curve

Table 8. Loss trajectory over the 45 epochs. Source: `reports/baseline_resnet50/history_fold0.csv`.

PHASE	EPOCHS	TRAINING LOSS	VALIDATION LOSS
Head only (backbone frozen)	0–4	1.42 → 0.83	1.80 → 1.71
Full fine-tune	5–44	0.82 → 0.49	1.68 → 0.97
Whole run	45	1.42 → 0.49	1.80 → 0.97

Figure 8. Training curves, adopted ResNet-50 run (45 epochs)

Validation loss is still falling at the end — the run stopped on the epoch budget, not on early stopping




Figure 2. Both curves fall monotonically and neither turns upward, so this run shows **no overfitting signature** — the gap between them is roughly constant after epoch 20 rather than widening. Note that validation loss was *still decreasing* at epoch 44: the run ended because it hit its epoch budget, not because it converged. The absolute values are inflated by label smoothing, which puts a floor under the cross-entropy, so they should be read as a trend and not compared against runs trained without it.

IV. Backend Deployment and Tech Stack

IV.1 Backend service framework

Flask 3.1.3 (`website/app.py`), serving both the API and the frontend from one process. It calls the repository's real inference code — `src.model`, `src.data`, `src.evaluate` — rather than a reimplementaion, so what the site serves is exactly what was evaluated.

COMPONENT	CHOICE
Framework	Flask 3.1.3, built-in development server, <code>127.0.0.1:5001</code> by default
Inference	PyTorch 2.13.0 + timm 1.0.28; device selectable <code>cpu</code> / <code>cuda</code> / <code>mps</code> , default CPU
Endpoints	<code>POST /api/identify</code> — multipart file or base64 JSON, returns class, calibrated confidence, all six probabilities, an OOD flag and an attention bounding box. <code>GET /api/model</code> — served checkpoint, image size, device, calibration state.
Frontend	Static <code>index.html</code> + <code>support.js</code> , React 18 UMD vendored locally under <code>website/vendor/</code> so the site runs with no network access
Model loading	Once at startup; accepts a single <code>best.pth</code> or an <code>ensemble.json</code> manifest

IV.2 Measured inference and API response time

The figures below were **measured, not estimated**: 12 real HTTP requests per mode against the running server after warm-up, plus 30 direct model calls, on the development machine (Apple Silicon, PyTorch 2.13.0, 4 threads).

Table 9. End-to-end `POST /api/identify` latency, including HTTP, JPEG decode, preprocessing, inference, calibration, the out-of-distribution check and the attention box.

MODE	MEDIAN	MIN	MAX	NOTES
Standard (1 view)	46 ms	44 ms	50 ms	Default path
Accuracy mode (4-view TTA)	120 ms	113 ms	137 ms	~2.6× the work

Table 10. Model-only latency (preprocessing + forward pass + softmax), excluding HTTP and the OOD/attention extras.

DEVICE	SINGLE VIEW	4-VIEW TTA
CPU (deployment default)	19.4 ms	60.2 ms
Apple GPU (MPS)	8.2 ms	24.3 ms

The difference between Tables 9 and 10 — roughly 26 ms — is the honest cost of everything that is not the neural network: HTTP handling, JPEG decoding, the feature-space out-of-distribution comparison against a 2,166-embedding bank, and computing the class-activation bounding box.

INTERPRETATION

Both modes are comfortably interactive; the standard path is well under the ~100 ms threshold at which a user perceives a response as instantaneous. Two caveats for honesty: this is the Flask **development** server, which handles one request at a time and is not a production deployment, and these timings are from the development machine rather than the presentation laptop.

V. Difficulties, Improvements, and Future Work

V.1 The main difficulty: a benchmark hidden inside another public dataset

While expanding the corpus we discovered that `omasteam/waste-garbage-management-dataset` silently **contains TrashNet**. Because our downloader collapses exact pixel duplicates, 2,218 of the 2,527 TrashNet images ended up stored under `garbage_classification_*.jpg` filenames — including **all 361 images of our already-spent test set**, 318 of them under the disguised names. Training on the corpus as downloaded would have meant training on 100% of the test set.

It was recoverable exactly. The import routine walked `sorted(zip members)`, so each image's index is a deterministic function of its original filename; reconstructing that ordering and joining on the index matched **2,527 / 2,527** rows with complete label agreement. Every affected near-duplicate group is now forced into the test split, and an assertion re-checks it on every load.

V.2 Other difficulties and what they taught us

DIFFICULTY	RESOLUTION AND LESSON
Accuracy improved while the minority class got worse	Modern augmentation gained 3 images overall and lost 2 of 23 <i>trash</i> images. Per-class recall became part of the acceptance gate, not a post-hoc observation.
A guard that guarded nothing	Our documentation claimed a file "physically blocks re-running the final evaluation". Searching the source found no code reading it. Single-use evaluation is now genuinely enforced. A documented guarantee is not an implemented one.
Ignore rules that ignored nothing	<code>.gitignore</code> patterns were root-anchored and never matched the real directories, so files we believed were deliberately excluded simply were not. Verified afterwards with <code>git check-ignore</code> .
A test whose name overstated it	A unit test named for a fallback never reached that code path. It passed, and its name was false. A test name is a claim about coverage.
Confidence was not correctness	Label smoothing left the model badly under-confident. The NLL-optimal temperature (0.22) made some <i>wrong</i> answers look 95% certain, so we chose a conservative 0.65 instead.
The closed-set problem	A six-class softmax must assign a photo of a cat to one of six classes. Margin, entropy and energy guards all failed — one real bottle looked more uncertain than the cat. Replaced with feature-space nearest-neighbour distance (threshold 0.70).
Scaling the duplicate search	Pairwise similarity built dense (N, N) matrices — fine at N = 2,527, about 0.52 GB each at N = 11,359. Rewritten to process in row-blocks.
Cross-platform paths	Four sites serialised paths in a way that would emit backslashes on Windows and break the manifest joins — on the exact machine where training was to run. Found by inspection before it failed.

V.3 Improvements already implemented

- **Corpus expansion** from 2,527 to 11,299 images across three domains, with the *trash* class rising from 137 to 1,033 — the minority-class problem addressed by data rather than by technique.
- **Parameter cap enforced in code**, so an ineligible backbone fails immediately instead of after hours of training.
- **Automatic GPU utilisation** on CUDA: TF32, cuDNN autotuning, channels-last memory format, and bfloat16 autocast on Ampere or newer. Enabling bfloat16 exposed a real defect — the loop keyed autocast off "a gradient scaler exists", which is wrong for bfloat16 and would have silently disabled mixed precision.
- **Batch-size autotuning** that probes real training steps until the GPU runs out of memory, applies the linear learning-rate scaling rule, and refuses to scale past 4×.

- **Preflight verification** that blocks training unless split arithmetic, disjointness, quarantine containment and model eligibility all pass.
- **Full resumability** — model, optimizer, scheduler, phase, EMA and early-stopping counter are all checkpointed, so an interrupted run loses at most one epoch.

V.4 The planned final run — a 13-hour budget

Remote GPU time is limited to roughly **13 hours**, which is the binding constraint on everything below. The plan is deliberately narrow: **one strong ConvNeXtV2 Tiny model**, trained well, rather than a broad comparison that produces four mediocre ones.

Table 11. Planned schedule. Every stage is resumable, so an interruption costs at most one epoch.

STAGE	BUDGET	PURPOSE
Setup, corpus download, split rebuild, preflight	30–90 min	Must exit 0 before anything trains
GPU batch-size autotune	5–15 min	Every later stage runs faster; pays for itself immediately
ConvNeXtV2 Tiny at 224px	2–4 h	The primary model. A usable checkpoint exists from here on.
Progressive fine-tune at 384px	3–6 h	Warm-started from the 224px checkpoint; historically the largest single gain
Spare capacity — one extra fold, or TTA and validation checks	remainder	Absorbs overruns; only used if the earlier stages finish early
Champion update, packaging, transfer	15–30 min	Reserved, not optional
Total	≈ 7–12 h	Leaves deliberate slack inside the 13-hour window

Deliberately skipped, and why. Each of these is methodologically sound and does not fit the budget:

SKIPPED	REASON
Four-model backbone bake-off	Would consume most of the budget selecting an architecture instead of training one properly. The three alternative configs stay in the repository for anyone with more time.
Full five-fold cross-validation	Five times the training cost. It will not fit safely, so results will be single-fold and must be reported as such.
Two-architecture ensemble	Requires the bake-off first. Also raises an interpretation question about the 30M cap that we would rather not depend on.
Additional datasets	Five sources are already configured and pinned; more would mean re-downloading and re-splitting.
Manual label-noise cleaning	Needs out-of-fold predictions across all five folds before the audit can even begin.

WHAT THIS MEANS FOR HOW RESULTS MUST BE REPORTED

Because five-fold training is being skipped, the Phase 2 number will be a **single-fold validation result**, not a mean with a standard deviation. It should be captioned that way. On a 1,912-image validation fold, a difference of under roughly 0.5 macro-F1 points is noise and should not be claimed as an improvement.

V.5 Longer-term work, beyond this budget

1. **Five-fold cross-validation and ensembling**, for a mean and standard deviation rather than a single fold, and for averaged predictions across folds.
2. **Label-noise cleaning.** An earlier audit found roughly half of the worst errors were mislabeled or genuinely ambiguous, and we have already proved that three image pairs carry contradictory labels across sources. Ranking out-

of-fold predictions by loss and auditing the worst ~ 300 would give a measurable before/after delta — a good result in its own right.

3. **Cross-source overlap audit.** We proved TrashNet's overlap only because its import ordering made exact recovery possible. We have *no* equivalent guarantee about the other four sources, and given what we found once, this is the most likely place for another surprise.
4. **A production server** (Gunicorn or Uvicorn) if the site is ever exposed beyond a laptop, and a proper out-of-distribution benchmark instead of a single cat photograph.

Provenance of these numbers. Accuracy, precision, recall, F1 and the loss trajectory are read from `experiments.csv` and the files under `model/reports/`. Split counts are recomputed from the committed CSVs. Hyperparameters are quoted from `configs/baseline.yaml` and the logged `config.json`. Latency was measured on 22 July 2026 against the running server. Figures are generated by `model/paper/make_figures.py`. Nothing in this document is a recalled or estimated value.